

# Bank Marketing Data

Group 8

## Load Data

```
#Read dataset  
bank_df <- read_delim("bank-additional-full.csv", delim=";")
```

```
## Warning in gzfile(file, mode): cannot open compressed file 'C:/Users/yohci/  
## AppData/Local/Temp/RtmpMtZd46\\file4eac5ef453ff', probable reason 'No such file  
## or directory'
```

```
##  
## -- Column specification -----  
## cols(  
##   .default = col_character(),  
##   age = col_double(),  
##   duration = col_double(),  
##   campaign = col_double(),  
##   pdays = col_double(),  
##   previous = col_double(),  
##   emp.var.rate = col_double(),  
##   cons.price.idx = col_double(),  
##   cons.conf.idx = col_double(),  
##   euribor3m = col_double(),  
##   nr.employed = col_double()  
## )  
## i Use `spec()` for the full column specifications.
```

```

#Assign category to all categorical variables

#2.job as category
bank_df$job <- as.factor(bank_df$job)

#3.marital status as category
bank_df$marital <- as.factor(bank_df$marital)

#4.education as category
bank_df$education <- as.factor(bank_df$education)

#5.credit default as category
bank_df$default <- as.factor(bank_df$default)

#6.housing loan as category
bank_df$housing <- as.factor(bank_df$housing)

#7.personal loan as category
bank_df$loan <- as.factor(bank_df$loan)

#8.contact communication type as category
bank_df$contact <- as.factor(bank_df$contact)

#9.last contact month of year as category
bank_df$month <- as.factor(bank_df$month)

#10.last contact day of the month as category
bank_df$day_of_week <- as.factor(bank_df$day_of_week)

#15.outcome of the previous marketing campaign as category
bank_df$poutcome <- as.factor(bank_df$poutcome)

#21.output y as binary factor
bank_df$y <- factor(bank_df$y, levels = c("no","yes"))

dim(bank_df)

```

```
## [1] 41188    21
```

## Data preprocessing

```

bank_df %>%
  summarise_all(list(~sum(. == "unknown"))) %>%
  gather(key = "variable", value = "nr_unknown") %>%
  arrange(-nr_unknown)

```

```
## # A tibble: 21 x 2
##   variable    nr_unknown
##   <chr>      <int>
## 1 default      8597
## 2 education    1731
## 3 housing       990
## 4 loan          990
## 5 job           330
## 6 marital       80
## 7 age           0
## 8 contact       0
## 9 month         0
## 10 day_of_week  0
## # ... with 11 more rows
```

```
# Analyse default
table(bank_df$default)
```

```
##
##      no unknown    yes
## 32588    8597      3
```

```
## This is not usable, too few "yes" to evaluate
```

## analyse the unknown values

```
# setting default parameters for crosstables
# fun_crosstable = function(df, var1, var2){
#   # df: dataframe containing both columns to cross
#   # var1, var2: columns to cross together.
#   CrossTable(df$var1, df$var2,
#     prop.r = T,
#     prop.c = F,
#     prop.t = F,
#     prop.chisq = F,
#     dnn = c(var1, var2)) # dimension names
# }

#default
CrossTable(bank_df$default, bank_df$y, prop.r = T, prop.c=F, prop.chisq=F, dnn = c("default",
"y"))
```

```
##
##
##   Cell Contents
## |-----|
## |                N |
## |      N / Row Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  41188
##
##
##           | y
##   default |    no |    yes | Row Total |
## -----|-----|-----|-----|
##         no |  28391 |   4197 |   32588 |
##           |    0.871 |   0.129 |   0.791 |
##           |    0.689 |   0.102 |         |
## -----|-----|-----|-----|
##        unknown |   8154 |    443 |    8597 |
##           |    0.948 |   0.052 |   0.209 |
##           |    0.198 |   0.011 |         |
## -----|-----|-----|-----|
##         yes |      3 |      0 |      3 |
##           |    1.000 |   0.000 |   0.000 |
##           |    0.000 |   0.000 |         |
## -----|-----|-----|-----|
## Column Total |  36548 |   4640 |   41188 |
## -----|-----|-----|-----|
##
##
```

```
table(bank_df$default)
```

```
##
##      no unknown    yes
##  32588    8597      3
```

```
# job
CrossTable(bank_df$job, bank_df$y, prop.r = T, prop.c=F, prop.chisq=F, dnn = c("job", "y"))
```

```

##
##
##      Cell Contents
## |-----|
## |                N |
## |      N / Row Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  41188
##
##
##      | y
##      job |      no |      yes | Row Total |
## -----|-----|-----|-----|
##      admin. |      9070 |      1352 |      10422 |
##              |      0.870 |      0.130 |      0.253 |
##              |      0.220 |      0.033 |              |
## -----|-----|-----|-----|
##      blue-collar |      8616 |      638 |      9254 |
##              |      0.931 |      0.069 |      0.225 |
##              |      0.209 |      0.015 |              |
## -----|-----|-----|-----|
##      entrepreneur |      1332 |      124 |      1456 |
##              |      0.915 |      0.085 |      0.035 |
##              |      0.032 |      0.003 |              |
## -----|-----|-----|-----|
##      housemaid |      954 |      106 |      1060 |
##              |      0.900 |      0.100 |      0.026 |
##              |      0.023 |      0.003 |              |
## -----|-----|-----|-----|
##      management |      2596 |      328 |      2924 |
##              |      0.888 |      0.112 |      0.071 |
##              |      0.063 |      0.008 |              |
## -----|-----|-----|-----|
##      retired |      1286 |      434 |      1720 |
##              |      0.748 |      0.252 |      0.042 |
##              |      0.031 |      0.011 |              |
## -----|-----|-----|-----|
##      self-employed |      1272 |      149 |      1421 |
##              |      0.895 |      0.105 |      0.035 |
##              |      0.031 |      0.004 |              |
## -----|-----|-----|-----|
##      services |      3646 |      323 |      3969 |
##              |      0.919 |      0.081 |      0.096 |
##              |      0.089 |      0.008 |              |
## -----|-----|-----|-----|
##      student |      600 |      275 |      875 |
##              |      0.686 |      0.314 |      0.021 |
##              |      0.015 |      0.007 |              |
## -----|-----|-----|-----|
##      technician |      6013 |      730 |      6743 |
##              |      0.892 |      0.108 |      0.164 |
##              |      0.146 |      0.018 |              |
## -----|-----|-----|-----|
##      unemployed |      870 |      144 |      1014 |


```

```
##           |      0.858 |      0.142 |      0.025 |
##           |      0.021 |      0.003 |              |
## -----|-----|-----|-----|
##      unknown |      293 |      37 |      330 |
##           |      0.888 |      0.112 |      0.008 |
##           |      0.007 |      0.001 |              |
## -----|-----|-----|-----|
## Column Total |     36548 |      4640 |     41188 |
## -----|-----|-----|-----|
##
##
```

```
table(bank_df$job)
```

```
##
##      admin.  blue-collar  entrepreneur  housemaid  management
##      10422      9254      1456      1060      2924
##      retired self-employed  services  student  technician
##      1720      1421      3969      875      6743
##      unemployed  unknown
##      1014      330
```

```
bank_df %>%
  ggplot() +
  geom_mosaic(aes(x = product(y, job), fill = y)) +
  #mosaic_theme +
  xlab("Job") +
  ylab(NULL)
```



```
bank_df <- bank_df %>%
  mutate(job = recode(job, "unknown" = "unconventional"))

# marital
CrossTable(bank_df$marital, bank_df$y, prop.r = T, prop.c=F, prop.chisq=F, dnn = c("marital",
"y"))
```


```
##
##
##   Cell Contents
## |-----|
## |                N |
## |      N / Row Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  41188
##
##
```

	y		
marital	no	yes	Row Total
divorced	4136	476	4612
	0.897	0.103	0.112
	0.100	0.012	
married	22396	2532	24928
	0.898	0.102	0.605
	0.544	0.061	
single	9948	1620	11568
	0.860	0.140	0.281
	0.242	0.039	
unknown	68	12	80
	0.850	0.150	0.002
	0.002	0.000	
Column Total	36548	4640	41188

```
##
##
```

```
## can merge single+unknown, married+divorced since values are similar
bank_df = bank_df %>%
  mutate(marital = recode(marital, "unknown" = "single", "divorced"="married"))

# education
bank_df %>%
  ggplot() +
  geom_mosaic(aes(x = product(y, education), fill = y)) +
  #mosaic_theme +
  xlab("Education") +
  ylab(NULL)
```




```
## recode unknown as univeristy degree because proportions are similar
bank_df = bank_df %>%
  mutate(education = recode(education, "unknown" = "university.degree"))

# housing
CrossTable(bank_df$housing, bank_df$y, prop.r = T, prop.c=F, prop.chisq=F, dnn = c("housing",
"y"))
```



```
##
##
##   Cell Contents
## |-----|
## |                N |
## |      N / Row Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  41188
##
##
##           | y
##   housing |    no |    yes | Row Total |
## -----|-----|-----|-----|
##         no |  16596 |   2026 |   18622 |
##           |    0.891 |   0.109 |   0.452 |
##           |    0.403 |   0.049 |         |
## -----|-----|-----|-----|
##        unknown |    883 |    107 |    990 |
##           |    0.892 |   0.108 |   0.024 |
##           |    0.021 |   0.003 |         |
## -----|-----|-----|-----|
##         yes |  19069 |   2507 |   21576 |
##           |    0.884 |   0.116 |   0.524 |
##           |    0.463 |   0.061 |         |
## -----|-----|-----|-----|
## Column Total |  36548 |   4640 |   41188 |
## -----|-----|-----|-----|
##
##
```

```
bank_df %>%
  ggplot() +
  geom_mosaic(aes(x = product(y, housing), fill = y)) +
  #mosaic_theme +
  xlab("housing") +
  ylab(NULL)
```



```
## the plot looks very similar, do chisquared test to see if there are differences
chisq.test(bank_df$housing, bank_df$y) # drop this column
```

```
##
## Pearson's Chi-squared test
##
## data: bank_df$housing and bank_df$y
## X-squared = 5.6845, df = 2, p-value = 0.05829
```

```
bank_df$housing <- NULL

# Loan
chisq.test(bank_df$loan, bank_df$y) # drop col, pvalue >0.1
```


```
##
## Pearson's Chi-squared test
##
## data: bank_df$loan and bank_df$y
## X-squared = 1.094, df = 2, p-value = 0.5787
```

```

bank_df$loan <- NULL

# pdays
# poutcome
bank_df %>%
  ggplot() +
  geom_mosaic(aes(x = product(y, poutcome), fill = y)) +
  #mosaic_theme +
  xlab("poutcome") +
  ylab(NULL)

```



```

bank_df = bank_df %>%
  mutate(past_dummyvar = recode(poutcome, "failure" = 0.5, "nonexistent"=0.2, "success"=1))
# combining previous and poutcome
bank_df$past_dummyvar1 = bank_df$past_dummyvar*(bank_df$previous+1)
chisq.test(bank_df$past_dummyvar1, bank_df$y)

```

```

## Warning in chisq.test(bank_df$past_dummyvar1, bank_df$y): Chi-squared
## approximation may be incorrect

```

```

##
## Pearson's Chi-squared test
##
## data: bank_df$past_dummyvar1 and bank_df$y
## X-squared = 4383.4, df = 11, p-value < 2.2e-16

```

```

bank_df$previous <-NULL
bank_df$poutcome <-NULL
bank_df$past_dummyvar <-NULL

bank_df = bank_df %>%
  mutate(pdays_dummy = if_else(pdays == 999, "0", "1")) %>%
  select(-pdays)
bank_df$pdays<-NULL

#resolve default, let yes become unknown
bank_df = bank_df %>%
  mutate(default = recode(default, "yes"="unknown"))

# dayofweek
bank_df = bank_df %>%
  mutate(day_of_week = recode(day_of_week, "mon"=1, "tue"=2,"wed"=3,"thu"=4,"fri"=5))

# age
bank_df = bank_df %>%
  mutate(age = if_else(
    age<20, 1, if_else(
      age<23, 2, if_else(
        age<26, 3, if_else(
          age<31, 4, if_else(
            age<41, 5, if_else(age<51, 6, 7)))))))

#dataset after preprocessing
dim(bank_df)

```

```
## [1] 41188    18
```

```
summary(bank_df)
```

```
##          age          job          marital
## Min.   :1.000   admin.   :10422   married:29540
## 1st Qu.:5.000   blue-collar: 9254   single :11648
## Median :5.000   technician : 6743
## Mean   :5.367   services   : 3969
## 3rd Qu.:6.000   management : 2924
## Max.    :7.000   retired    : 1720
##          (Other)   : 6156
##          education      default      contact      month
## basic.4y      : 4176   no      :32588   cellular :26144   may      :13769
## basic.6y      : 2292   unknown: 8600   telephone:15044   jul      : 7174
## basic.9y      : 6045
## high.school   : 9515
## illiterate    :   18
## professional.course: 5243
## university.degree :13899
##          (Other): 2016
##   day_of_week   duration      campaign      emp.var.rate
## Min.   :1.00   Min.    : 0.0   Min.    : 1.000   Min.    :-3.40000
## 1st Qu.:2.00   1st Qu.: 102.0   1st Qu.: 1.000   1st Qu.: -1.80000
## Median :3.00   Median : 180.0   Median : 2.000   Median : 1.10000
## Mean   :2.98   Mean    : 258.3   Mean    : 2.568   Mean    : 0.08189
## 3rd Qu.:4.00   3rd Qu.: 319.0   3rd Qu.: 3.000   3rd Qu.: 1.40000
## Max.    :5.00   Max.    :4918.0   Max.    :56.000   Max.    : 1.40000
##
## cons.price.idx cons.conf.idx   euribor3m   nr.employed   y
## Min.   :92.20   Min.    :-50.8   Min.    :0.634   Min.    :4964   no :36548
## 1st Qu.:93.08   1st Qu.: -42.7   1st Qu.:1.344   1st Qu.:5099   yes: 4640
## Median :93.75   Median : -41.8   Median :4.857   Median :5191
## Mean   :93.58   Mean    : -40.5   Mean    :3.621   Mean    :5167
## 3rd Qu.:93.99   3rd Qu.: -36.4   3rd Qu.:4.961   3rd Qu.:5228
## Max.    :94.77   Max.    : -26.9   Max.    :5.045   Max.    :5228
##
## past_dummyvar1   pdays_dummy
## Min.    :0.2000   Length:41188
## 1st Qu.:0.2000   Class :character
## Median :0.2000   Mode  :character
## Mean    :0.3703
## 3rd Qu.:0.2000
## Max.    :8.0000
##
```

## Oversampling

```
n <- nrow(bank_df); n
```

```
## [1] 41188
```

```
majorind <- (1:n)[bank_df$y == "no"]
minorind <- (1:n)[bank_df$y == "yes"]
majorn <- length(majorind)
minorn <- length(minorind)

#sample(data_index, numberofdata, replacement?)
OSind<-sample(minorind,majorn-minorn,replace=TRUE)
OSdata<-bank_df[OSind,] # Length 4244
# Get the new combined and scaled dataset
OSdata_combined <- rbind(bank_df, bank_df[OSind,]) # Length 9066 = 4822+4244
table(OSdata_combined$y) # 4533 points each
```

```
##
##      no    yes
## 36548 36548
```

```
# splitting train and test
library(caTools)
set.seed(1)
smp_size <- floor(0.8*nrow(OSdata_combined))
train_ind <- sample(seq_len(nrow(OSdata_combined)), size = smp_size)
train <- OSdata_combined[train_ind, ]
test <- OSdata_combined[-train_ind, ]
table(train$y)
```

```
##
##      no    yes
## 29269 29207
```

```
table(test$y)
```

```
##
##      no    yes
## 7279 7341
```

# KNN

```

set.seed(8)
#use K-fold CV to find best trCtrl:
trctrl <- trainControl(method = "repeatedcv", number = 5, repeats = 1) #5 fold CV repeated 1
times

# knn.fit <- train(y ~., data = train, method = "knn",
#                 trControl=trctrl, tuneLength = 10) # tuneLength parameter tells the algo
rithm to try different default values for the main parameter

knn.fit <- train(y ~., data = train, method = "knn",
                trControl=trctrl) # tuneLength parameter tells the algorithm to try differen
t default values for the main parameter

# knn.fit <- train(y ~., data = train, method = "knn")#by default bootstrap is used to find t
uning parameter -> trCtrl
knn.fit

```

```

## k-Nearest Neighbors
##
## 58476 samples
##    17 predictor
##    2 classes: 'no', 'yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 1 times)
## Summary of sample sizes: 46781, 46780, 46780, 46781, 46782
## Resampling results across tuning parameters:
##
##  k  Accuracy  Kappa
##  5  0.9011048  0.8022425
##  7  0.8927595  0.7855535
##  9  0.8885355  0.7771034
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 5.

```

```

#predict using test data
knn.pred <- predict(knn.fit, newdata = test)
#knn.pred

#confusion matrix
cm.knn <- table(knn.pred, test$y)
cm.knn

```

```

##
## knn.pred   no  yes
##         no 6046  93
##         yes 1233 7248

```

```
TP <- cm.knn[2,2]
TN <- cm.knn[1,1]
FP <- cm.knn[2,1]
FN <- cm.knn[1,2]
```

```
#FPR / Type I error
FPR.knn = FP/(FP+TN)
FPR.knn
```

```
## [1] 0.1693914
```

```
#FNR / Type II error
FNR.knn = FN/(FN+TP)
FNR.knn
```

```
## [1] 0.01266857
```

```
#Precision
precis.knn = TP/(TP+FP)
precis.knn
```

```
## [1] 0.8546162
```

```
#Recall / sensitivity
recall.knn = TP/(TP+FN)
recall.knn
```

```
## [1] 0.9873314
```

```
#misclassification error
test.err.knn = 1-(sum(diag(cm.knn))/sum(cm.knn))
test.err.knn
```

```
## [1] 0.09069767
```

## Logistic Regression

```
set.seed(8)
glm.fit <- glm(y ~., data = train, family = binomial)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(glm.fit)
```



```
##
## Call:
## glm(formula = y ~ ., family = binomial, data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -8.4904  -0.3793  -0.1126   0.4861   2.9677
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -2.721e+02  2.326e+01 -11.700 < 2e-16 ***
## age           -4.289e-02  1.448e-02  -2.962 0.003059 **
## jobblue-collar -2.391e-01  4.974e-02  -4.808 1.52e-06 ***
## jobentrepreneur -1.097e-01  7.658e-02  -1.433 0.151953
## jobhousemaid    9.750e-03  9.554e-02   0.102 0.918720
## jobmanagement  -6.804e-02  5.453e-02  -1.248 0.212142
## jobretired      5.638e-01  6.612e-02   8.526 < 2e-16 ***
## jobself-employed -2.345e-01  7.518e-02  -3.119 0.001817 **
## jobservices    -1.739e-01  5.406e-02  -3.217 0.001294 **
## jobstudent      2.936e-01  7.886e-02   3.723 0.000197 ***
## jobtechnician  -3.499e-02  4.619e-02  -0.758 0.448709
## jobunemployed   1.733e-01  8.409e-02   2.061 0.039328 *
## jobunconventional 2.087e-01  1.459e-01   1.430 0.152748
## maritalsingle   8.718e-02  3.282e-02   2.656 0.007904 **
## educationbasic.6y 1.427e-02  7.680e-02   0.186 0.852612
## educationbasic.9y -5.231e-02  6.047e-02  -0.865 0.387027
## educationhigh.school 1.075e-02  5.888e-02   0.183 0.855151
## educationilliterate 1.308e+00  5.038e-01   2.597 0.009411 **
## educationprofessional.course 1.366e-01  6.543e-02   2.088 0.036825 *
## educationuniversity.degree 2.348e-01  5.764e-02   4.073 4.64e-05 ***
## defaultunknown  -3.289e-01  4.085e-02  -8.052 8.15e-16 ***
## contacttelephone -4.912e-01  4.927e-02  -9.970 < 2e-16 ***
## monthaug        1.095e+00  8.526e-02  12.847 < 2e-16 ***
## monthdec        1.925e-01  1.624e-01   1.185 0.236050
## monthjul        -6.024e-03  6.205e-02  -0.097 0.922660
## monthjun        -9.603e-01  7.873e-02 -12.198 < 2e-16 ***
## monthmar        2.154e+00  9.995e-02  21.554 < 2e-16 ***
## monthmay        -7.900e-01  5.171e-02 -15.279 < 2e-16 ***
## monthnov        -6.764e-01  7.673e-02  -8.815 < 2e-16 ***
## monthoct        4.405e-01  1.009e-01   4.366 1.27e-05 ***
## monthsep        4.440e-01  1.170e-01   3.794 0.000148 ***
## day_of_week     -2.724e-03  9.377e-03  -0.290 0.771461
## duration        6.965e-03  6.679e-05 104.283 < 2e-16 ***
## campaign        -2.446e-02  6.991e-03  -3.498 0.000469 ***
## emp.var.rate    -2.325e+00  8.776e-02 -26.498 < 2e-16 ***
## cons.price.idx   2.557e+00  1.547e-01  16.528 < 2e-16 ***
## cons.conf.idx    3.846e-03  5.471e-03   0.703 0.482038
## euribor3m       6.296e-01  8.431e-02   7.468 8.13e-14 ***
## nr.employed     5.489e-03  1.908e-03   2.877 0.004019 **
## past_dummyvar1  -2.865e-01  3.438e-02  -8.332 < 2e-16 ***
## pdays_dummy1    2.157e+00  9.588e-02  22.501 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 81065  on 58475  degrees of freedom
```

```
## Residual deviance: 38198  on 58435  degrees of freedom
## AIC: 38280
##
## Number of Fisher Scoring iterations: 6
```

```
#predict using test data
glm.prob <- predict(glm.fit, type = "response", newdata = test)

#check which one is 'Yes'
contrasts(test$y)#Yes = 1, Low = 0
```

```
##      yes
## no      0
## yes     1
```

```
glm.pred <- rep('no', nrow(test))
glm.pred[glm.prob > 0.5] <- 'yes' #yes = 1, no = 0

#confusion matrix
cm.reg = table(glm.pred, test$y)
cm.reg
```

```
##
## glm.pred   no  yes
##          no 6274 849
##          yes 1005 6492
```

```
TP <- cm.reg[2,2]
TN <- cm.reg[1,1]
FP <- cm.reg[2,1]
FN <- cm.reg[1,2]

#FPR / Type I error
FPR.reg = FP/(FP+TN)
FPR.reg
```

```
## [1] 0.1380684
```

```
#FNR / Type II error
FNR.reg = FN/(FN+TP)
FNR.reg
```

```
## [1] 0.1156518
```

```
#Precision
precis.reg = TP/(TP+FP)
precis.reg
```

```
## [1] 0.8659464
```

```
#Recall / sensitivity
recall.reg = TP/(TP+FN)
recall.reg
```

```
## [1] 0.8843482
```

```
#misclassification error
test.err.reg = 1-(sum(diag(cm.reg))/sum(cm.reg))
test.err.reg
```

```
## [1] 0.1268126
```

## Decision Tree

```
train.tree <- data.frame(train)
test.tree <- data.frame(test)

#Decision Tree using rpart()
set.seed(8)
#use K-fold CV to find best trCtrl:
trctrl <- trainControl(method = "repeatedcv", number = 5, repeats = 1) #5 fold CV repeated 1
times


cls.tree1 = train(y ~ ., data=train.tree, method="rpart",
                  trControl=trctrl)# tuneLength parameter tells the algorithm to try differen
t default values for the main parameter

# cls.tree1 = train(y ~ ., data=train.tree, method="rpart")#by default bootstrap is used to f
ind tuning parameter -> trCtrl

cls.tree1
```

```
## CART
##
## 58476 samples
##    17 predictor
##    2 classes: 'no', 'yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 1 times)
## Summary of sample sizes: 46781, 46780, 46780, 46781, 46782
## Resampling results across tuning parameters:
##
##    cp          Accuracy    Kappa
## 0.07155819  0.8109814  0.6219926
## 0.08609238  0.7862363  0.5725275
## 0.45174102  0.6809321  0.3617189
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.07155819.
```

```
plot(cls.tree1$finalModel)
text(cls.tree1$finalModel)
```



```
#predict using test data
tree.pred1 <- predict(cls.tree1, newdata = test.tree)
#tree.pred1

#confusion matrix
cm.tree1 <- table(tree.pred1, test$y)
cm.tree1
```

```
##
## tree.pred1   no  yes
##           no 6441 1946
##           yes 838 5395
```

```
TP <- cm.tree1[2,2]
TN <- cm.tree1[1,1]
FP <- cm.tree1[2,1]
FN <- cm.tree1[1,2]

#FPR / Type I error
FPR.tree1 = FP/(FP+TN)
FPR.tree1
```

```
## [1] 0.1151257
```

```
#FNR / Type II error
FNR.tree1 = FN/(FN+TP)
FNR.tree1
```

```
## [1] 0.2650865
```

```
#Precision
precis.tree1 = TP/(TP+FP)
precis.tree1
```

```
## [1] 0.8655543
```

```
#Recall / sensitivity
recall.tree1 = TP/(TP+FN)
recall.tree1
```

```
## [1] 0.7349135
```

```
#misclassification error
test.err.tree1 = 1-(sum(diag(cm.tree1))/sum(cm.tree1))
test.err.tree1
```

```
## [1] 0.1904241
```

## Random Forest

```
#Random forest with 500 bootstrapped trees
#p = 16
sqrt(16) # ntree = 4
```

```
## [1] 4
```

```
set.seed(8)
rf.cls <- randomForest(y ~ .,
                      data = train,
                      mtry = 4,
                      ntree = 500,
                      importance = TRUE)
rf.cls
```



```
##
## Call:
##  randomForest(formula = y ~ ., data = train, mtry = 4, ntree = 500,      importance = TRUE)
##
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 4
##
##           OOB estimate of  error rate: 4.37%
## Confusion matrix:
##           no  yes  class.error
## no  26729  2540  0.0867812361
## yes   14 29193  0.0004793371
```

```
#ls(rf.cls)
importance(rf.cls)
```

	no	yes	MeanDecreaseAccuracy	MeanDecreaseGini
age	12.429526	153.40158	152.30127	681.2940
job	23.572414	200.25575	203.72387	1323.2002
marital	2.934999	129.92328	129.00311	236.0142
education	4.331264	209.88006	208.62744	825.6247
default	-7.766226	73.69850	73.11891	199.4638
contact	8.627092	57.53414	51.95514	247.0378
month	31.939679	54.66100	56.11903	1072.9966
day_of_week	24.404648	203.07014	192.57306	696.6359
duration	310.415999	362.55453	381.64679	11950.4967
campaign	-4.966380	208.09632	206.54586	764.4524
emp.var.rate	20.019480	24.03849	27.90471	1272.2180
cons.price.idx	19.580034	42.07286	39.21067	533.6724
cons.conf.idx	18.503076	45.29669	40.77544	744.0300
euribor3m	28.065364	89.57287	76.13213	3135.1825
nr.employed	22.469122	38.78348	39.93334	2446.3922
past_dummyvar1	8.633164	63.41333	58.71083	668.3224
pdays_dummy	5.335515	86.15948	77.43937	467.9722

```
varImpPlot(rf.cls)
```

rf.cls



```
#predict using test data
rf.pred <- predict(rf.cls, newdata = test, type = "class")
#rf.pred
```

```
#confusion matrix
cm.rf <- table(rf.pred, test$y)
cm.rf
```

```
##
## rf.pred   no  yes
##      no 6677   1
##      yes  602 7340
```

```
TP <- cm.rf[2,2]
TN <- cm.rf[1,1]
FP <- cm.rf[2,1]
FN <- cm.rf[1,2]

#FPR / Type I error
FPR.rf = FP/(FP+TN)
FPR.rf
```

```
## [1] 0.08270367
```

```
#FNR / Type II error
```

```
FNR.rf = FN/(FN+TP)
```

```
FNR.rf
```

```
## [1] 0.0001362212
```

```
#Precision
```

```
precis.rf = TP/(TP+FP)
```

```
precis.rf
```

```
## [1] 0.9242005
```

```
#Recall / sensitivity
```

```
recall.rf = TP/(TP+FN)
```

```
recall.rf
```

```
## [1] 0.9998638
```

```
#misclassification error
```

```
test.err.rf = 1-(sum(diag(cm.rf))/sum(cm.rf))
```

```
test.err.rf
```

```
## [1] 0.04124487
```

## Gradient Boosting

```
#Gradient boosting
```

```
set.seed(8)
```

```
#Use K-fold CV to find best trControl
```

```
fitControl <- trainControl(method = "repeatedcv",
```

```
                           number = 5,
```

```
                           repeats = 1) #5 folds repeated 1 times
```

```
gbm.fit <- train(y ~ ., data = train,
```

```
                method = "gbm",
```

```
                trControl = fitControl,
```

```
                verbose = FALSE)
```

```
# gbm.fit <- train(y ~ ., data = train,
```

```
#                method = "gbm",
```

```
#                verbose = FALSE) #by default bootstrap is used to find tuning parameter ->
```

```
trCtrl
```

```
gbm.fit
```



```
## Stochastic Gradient Boosting
##
## 58476 samples
##    17 predictor
##    2 classes: 'no', 'yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 1 times)
## Summary of sample sizes: 46781, 46780, 46780, 46781, 46782
## Resampling results across tuning parameters:
##
##  interaction.depth  n.trees  Accuracy   Kappa
##  1                   50      0.8572065  0.7144247
##  1                   100     0.8711268  0.7422659
##  1                   150     0.8738971  0.7478057
##  2                    50     0.8727513  0.7455228
##  2                   100     0.8804125  0.7608453
##  2                   150     0.8846365  0.7692931
##  3                    50     0.8796088  0.7592395
##  3                   100     0.8853205  0.7706612
##  3                   150     0.8886039  0.7772268
##
## Tuning parameter 'shrinkage' was held constant at a value of 0.1
##
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were n.trees = 150, interaction.depth =
## 3, shrinkage = 0.1 and n.minobsinnode = 10.
```

```
#predict using test data
gbm.pred <- predict(gbm.fit, newdata = test)
#gbm.pred

#confusion matrix
cm.gbm <- table(gbm.pred, test$y)
cm.gbm
```

```
##
## gbm.pred   no  yes
##        no 6182 522
##        yes 1097 6819
```

```
TP <- cm.gbm[2,2]
TN <- cm.gbm[1,1]
FP <- cm.gbm[2,1]
FN <- cm.gbm[1,2]

#FPR / Type I error
FPR.gbm = FP/(FP+TN)
FPR.gbm
```

```
## [1] 0.1507075
```

```
#FNR / Type II error
FNR.gbm = FN/(FN+TP)
FNR.gbm
```

```
## [1] 0.07110748
```

```
#Precision
precis.gbm = TP/(TP+FP)
precis.gbm
```

```
## [1] 0.8614199
```

```
#Recall / sensitivity
recall.gbm = TP/(TP+FN)
recall.gbm
```

```
## [1] 0.9288925
```

```
#misclassification error
test.err.gbm = 1-(sum(diag(cm.gbm))/sum(cm.gbm))
test.err.gbm
```

```
## [1] 0.1107387
```

# AdaBoost

```
#AdaBoost
set.seed(8)
x.trainA = model.matrix(data=train, y~.-1)
y.trainA = rep(1, nrow(train))
y.trainA [train$y=="no"]=-1 #for Adaboost

x.testA = model.matrix(data=test, y~.-1)
y.testA = rep(1, nrow(test))
y.testA [test$y=="no"]=-1 #for Adaboost

ada.cls <- adaboost(x.trainA, y.trainA, tree_depth=5, n_rounds=500)
ada.cls
```

```
## AdaBoost: tree_depth = 5 rounds = 500
##
##
## In-sample confusion matrix:
## yhat
## y      -1      1
## -1 25704  3565
##  1  1187 28020
```

```
#predict using test data
ada.pred <- predict(ada.cls, x.testA)
#ada.pred

#confusion matrix
cm.ada <- table(ada.pred, y.testA) #-1 is "no", 1 is "yes"
cm.ada
```

```
##          y.testA
## ada.pred  -1    1
##          -1 6373 350
##           1  906 6991
```

```
TP <- cm.ada[2,2]
TN <- cm.ada[1,1]
FP <- cm.ada[2,1]
FN <- cm.ada[1,2]

#FPR / Type I error
FPR.ada = FP/(FP+TN)
FPR.ada
```

```
## [1] 0.1244676
```

```
#FNR / Type II error
FNR.ada = FN/(FN+TP)
FNR.ada
```

```
## [1] 0.04767743
```

```
#Precision
precis.ada = TP/(TP+FP)
precis.ada
```

```
## [1] 0.8852729
```

```
#Recall / sensitivity
recall.ada = TP/(TP+FN)
recall.ada
```

```
## [1] 0.9523226
```

```
#misclassification error
test.err.ada = 1-(sum(diag(cm.ada))/sum(cm.ada))
test.err.ada
```

```
## [1] 0.08590971
```

# XGBoost

```
#XGBoost
set.seed(8)
x.trainXG =model.matrix(data=train,y~.-1)
y.trainXG = rep(1, nrow(train))
y.trainXG[train$y=="no"]=0 #for XGBoost

x.testXG = model.matrix(data=test, y~.-1)
y.testXG = rep(1, nrow(test))
y.testXG[test$y=="no"]=0 #for XGBoost

xgb.cls <- xgboost(data=x.trainXG,label=y.trainXG,max_depth=5,eta=0.01,nrounds=500,verbose=FALSE)
#xgb.cls <- xgboost(data=x.trainXG,label=y.trainXG,max_depth=10,nrounds=500,verbose=FALSE)
xgb.cls
```

```
## ##### xgb.Booster
## raw: 1.1 Mb
## call:
##   xgb.train(params = params, data = dtrain, nrounds = nrounds,
##   watchlist = watchlist, verbose = verbose, print_every_n = print_every_n,
##   early_stopping_rounds = early_stopping_rounds, maximize = maximize,
##   save_period = save_period, save_name = save_name, xgb_model = xgb_model,
##   callbacks = callbacks, max_depth = 5, eta = 0.01)
## params (as set within xgb.train):
##   max_depth = "5", eta = "0.01", validate_parameters = "1"
## xgb.attributes:
##   niter
## callbacks:
##   cb.evaluation.log()
## # of features: 41
## niter: 500
## nfeatures : 41
## evaluation_log:
##   iter train_rmse
##     1    0.496814
##     2    0.493676
## ---
##    499    0.272070
##    500    0.272022
```

```
#predict using test data
xgb.pred.prob<-predict(xgb.cls,x.testXG)

xgb.pred<-as.numeric(xgb.pred.prob>0.5) #convert to 0 ("no") or 1 ("yes")

#confusion matrix
cm.xgb<-table(xgb.pred,y.testXG) #0 is "no", 1 is "yes"
cm.xgb
```

```
##          y.testXG
## xgb.pred    0    1
##          0 6226  391
##          1 1053 6950
```

```
TP <- cm.xgb[2,2]
TN <- cm.xgb[1,1]
FP <- cm.xgb[2,1]
FN <- cm.xgb[1,2]

#FPR / Type I error
FPR.xgb = FP/(FP+TN)
FPR.xgb
```

```
## [1] 0.1446627
```

```
#FNR / Type II error
FNR.xgb = FN/(FN+TP)
FNR.xgb
```

```
## [1] 0.0532625
```

```
#Precision
precis.xgb = TP/(TP+FP)
precis.xgb
```

```
## [1] 0.8684243
```

```
#Recall / sensitivity
recall.xgb = TP/(TP+FN)
recall.xgb
```

```
## [1] 0.9467375
```

```
#misclassification error
test.err.xgb = 1-sum(diag(cm.xgb))/sum(cm.xgb)
test.err.xgb
```

```
## [1] 0.09876881
```

## SVM with linear kernel

```

set.seed(8)
svm.fit <- svm(y~., data=train, kernel='linear', cost=1)
#summary(svm.fit)

#CV for tuning the cost parameter
set.seed(8)
tune.out1 <- tune(svm, y~.,
                  data=train,
                  kernel="linear",
                  )

#tune.out1 <- tune(svm, y~.,
#                  data=train,
#                  kernel="linear",
#                  ranges=list(cost=c(0.01, 0.1, 1, 10, 100)), tunecontrol=tune.control(cross=10))
summary(tune.out1)

```

```

##
## Error estimation of 'svm' using 10-fold cross validation: 0.1225287

```

```

svm.lin.best <- tune.out1$best.model
summary(svm.lin.best)

```

```

##
## Call:
## best.tune(method = svm, train.x = y ~ ., data = train, kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##         cost:  1
##
## Number of Support Vectors:  18932
##
## ( 9447 9485 )
##
##
## Number of Classes:  2
##
## Levels:
##   no yes

```

```

#predict using test data
lin.pred <- predict(svm.lin.best, test)

#confusion matrix
cm.lin <- table(lin.pred, test$y)
cm.lin

```

```
##
## lin.pred   no  yes
##          no 6149 639
##          yes 1130 6702
```

```
TP <- cm.lin[2,2]
TN <- cm.lin[1,1]
FP <- cm.lin[2,1]
FN <- cm.lin[1,2]

#FPR / Type I error
FPR.lin = FP/(FP+TN)
FPR.lin
```

```
## [1] 0.1552411
```

```
#FNR / Type II error
FNR.lin = FN/(FN+TP)
FNR.lin
```

```
## [1] 0.08704536
```

```
#Precision
precis.lin = TP/(TP+FP)
precis.lin
```

```
## [1] 0.8557201
```

```
#Recall / sensitivity
recall.lin = TP/(TP+FN)
recall.lin
```

```
## [1] 0.9129546
```

```
#misclassification error
test.err.lin = 1-(sum(diag(cm.lin))/sum(cm.lin))
test.err.lin
```

```
## [1] 0.1209986
```

## SVM with polynomial kernel

```

set.seed(8)
tune.out2 <- tune(svm, y~.,
                  data=train,
                  kernel="polynomial",
                  )

#tune.out2 <- tune(svm, y~.,
#                  data=train,
#                  kernel="polynomial",
#                  ranges=list(cost=c(0.1,1,5,10,15,20,50,100),
#                              degree=c(2,3,4)))
summary(tune.out2)

```

```

##
## Error estimation of 'svm' using 10-fold cross validation: 0.124085

```

```

svm.poly.best <- tune.out2$best.model
summary(svm.poly.best)

```

```

##
## Call:
## best.tune(method = svm, train.x = y ~ ., data = train, kernel = "polynomial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##     cost:    1
##    degree:   3
##   coef.0:   0
##
## Number of Support Vectors:  24284
##
## ( 12052 12232 )
##
##
## Number of Classes:  2
##
## Levels:
## no yes

```

```

#predict using test data
poly.pred <- predict(svm.poly.best, test)

#confusion matrix
cm.poly <- table(poly.pred, test$y)
cm.poly

```

```

##
## poly.pred   no  yes
##           no 6150 685
##           yes 1129 6656

```



```
TP <- cm.poly[2,2]
TN <- cm.poly[1,1]
FP <- cm.poly[2,1]
FN <- cm.poly[1,2]

#FPR / Type I error
FPR.poly = FP/(FP+TN)
FPR.poly
```

```
## [1] 0.1551037
```

```
#FNR / Type II error
FNR.poly = FN/(FN+TP)
FNR.poly
```

```
## [1] 0.09331154
```

```
#Precision
precis.poly = TP/(TP+FP)
precis.poly
```

```
## [1] 0.8549775
```

```
#Recall / sensitivity
recall.poly = TP/(TP+FN)
recall.poly
```

```
## [1] 0.9066885
```

```
#misclassification error
test.err.poly = 1-(sum(diag(cm.poly))/sum(cm.poly))
test.err.poly
```

```
## [1] 0.1240766
```

## SVM with rbf kernel

```
set.seed(8)
tune.out3 <- tune(svm, y~.,
                  data=train,
                  kernel="radial",)

#tune.out3 <- tune(svm, y~.,
#                  data=train,
#                  kernel="radial",
#                  ranges=list(cost=c(0.1,1,5,10),
#                               gamma=c(0.01,0.1,1,5,10,100)))
summary(tune.out3)
```

```
##
## Error estimation of 'svm' using 10-fold cross validation: 0.1132942
```

```
svm.rbf.best <- tune.out3$best.model
summary(svm.rbf.best)
```

```
##
## Call:
## best.tune(method = svm, train.x = y ~ ., data = train, kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##         cost:  1
##
## Number of Support Vectors:  17897
##
## ( 8758 9139 )
##
##
## Number of Classes:  2
##
## Levels:
##   no yes
```

```
#predict using test data
rbf.pred <- predict(svm.rbf.best, test)

#confusion matrix
cm.rbf <- table(rbf.pred, test$y)
cm.rbf
```

```
##
## rbf.pred   no  yes
##         no 6090 463
##         yes 1189 6878
```

```
TP <- cm.rbf[2,2]
TN <- cm.rbf[1,1]
FP <- cm.rbf[2,1]
FN <- cm.rbf[1,2]

#FPR / Type I error
FPR.rbf = FP/(FP+TN)
FPR.rbf
```

```
## [1] 0.1633466
```

```
#FNR / Type II error
FNR.rbf = FN/(FN+TP)
FNR.rbf
```

```
## [1] 0.06307043
```

```
#Precision
precis.rbf = TP/(TP+FP)
precis.rbf
```

```
## [1] 0.8526094
```

```
#Recall / sensitivity
recall.rbf = TP/(TP+FN)
recall.rbf
```

```
## [1] 0.9369296
```

```
#misclassification error
test.err.rbf = 1-(sum(diag(cm.rbf))/sum(cm.rbf))
test.err.rbf
```

```
## [1] 0.1129959
```

## Result Summary

```
options(digits = 3)
cl.err <- matrix(c(test.err.knn,FNR.knn,precis.knn,recall.knn,
                  test.err.reg,FNR.reg,precis.reg,recall.reg,
                  test.err.tree1,FNR.tree1,precis.tree1,recall.tree1,
                  test.err.rf,FNR.rf,precis.rf,recall.rf,
                  test.err.gbm,FNR.gbm,precis.gbm,recall.gbm,
                  test.err.ada,FNR.ada,precis.ada,recall.ada,
                  test.err.xgb,FNR.xgb,precis.xgb,recall.xgb,
                  test.err.lin,FNR.lin,precis.lin,recall.lin,
                  test.err.poly,FNR.poly,precis.poly,recall.poly,
                  test.err.rbf,FNR.rbf,precis.rbf,recall.rbf),
                ncol=4, byrow=TRUE)
colnames(cl.err) <- c('misclass error','type-II error','precision','recall')
rownames(cl.err) <- c('KNN',
                    'Logistic regression',
                    'Decision tree with rpart',
                    'Random forest',
                    'Gradient boosting',
                    'Adaboost',
                    'XGBoost',
                    'SVM with linear kernel',
                    'SVM with polynomial kernel',
                    'SVM with radial kernel')
as.table(cl.err)
```

##	misclass error	type-II error	precision	recall
## KNN	0.090698	0.012669	0.854616	0.987331
## Logistic regression	0.126813	0.115652	0.865946	0.884348
## Decision tree with rpart	0.190424	0.265087	0.865554	0.734913
## Random forest	0.041245	0.000136	0.924200	0.999864
## Gradient boosting	0.110739	0.071107	0.861420	0.928893
## Adaboost	0.085910	0.047677	0.885273	0.952323
## XGBoost	0.098769	0.053262	0.868424	0.946738
## SVM with linear kernel	0.120999	0.087045	0.855720	0.912955
## SVM with polynomial kernel	0.124077	0.093312	0.854978	0.906688
## SVM with radial kernel	0.112996	0.063070	0.852609	0.936930

Based on Type-II error comparison, best model is: Random Forest.

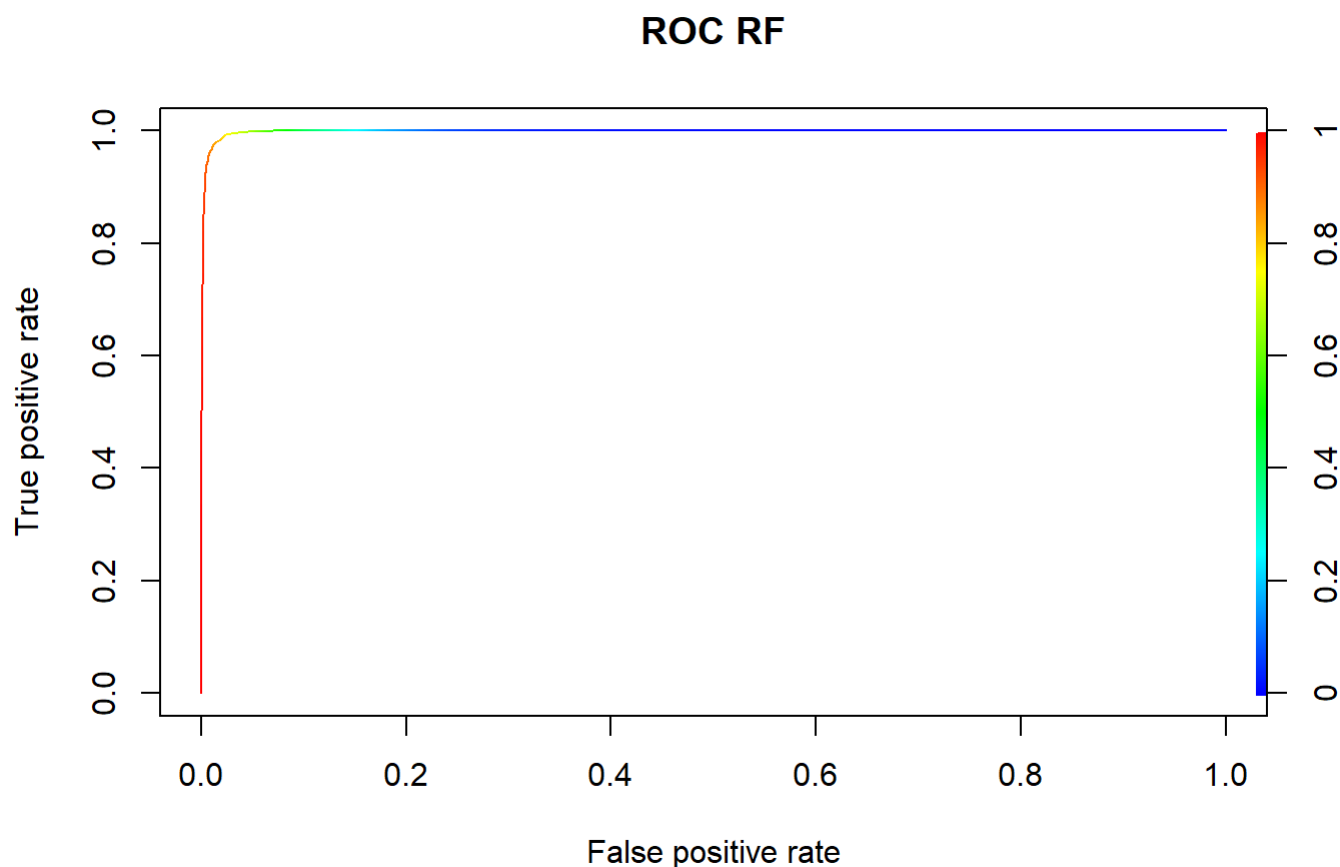
# ROC and AUC

## Random Forest

```
#Prepare model for ROC curve
rf.pred <- predict(rf.cls, newdata = test, type = "prob")

forestpred = prediction(rf.pred[,2], test$y)

roc.perf.rf = performance(forestpred, measure = "tpr", x.measure = "fpr")
plot(roc.perf.rf, main='ROC RF', colorize=T)
```



```

#Optimal cutoff
opt.cut = function(perf, pred){
  cut.ind = mapply(FUN=function(x, y, p){
    d = (x - 0)^2 + (y-1)^2
    ind = which(d == min(d))
    c(sensitivity = y[[ind]], specificity = 1-x[[ind]],
      cutoff = p[[ind]])
  }, perf@x.values, perf@y.values, pred@cutoffs)
}
#print(opt.cut(roc.perf.rf, forestpred))
roc.result = as.data.frame((opt.cut(roc.perf.rf, forestpred)))
roc.result

```

```

##              V1
## sensitivity 0.990
## specificity 0.979
## cutoff      0.784

```

```

rf.sens = roc.result[1,]
rf.spec = roc.result[2,]
rf.cutoff = roc.result[3,]

auc.perf.rf = performance(forestpred, measure = 'auc')
auc.rf = auc.perf.rf@y.values
auc.rf

```

```

## [[1]]
## [1] 0.998

```

## KNN

```


#Prepare model for ROC curve
knn.pred <- predict (knn.fit, test, type = "prob")

knnpred = prediction(knn.pred[,2], test$y)

roc.perf.knn = performance(knnpred, measure = "tpr", x.measure = "fpr")
plot(roc.perf.knn, main='ROC KNN', colorize=T)

```

## ROC KNN



```
#Optimal cutoff
opt.cut = function(perf, pred){
  cut.ind = mapply(FUN=function(x, y, p){
    d = (x - 0)^2 + (y-1)^2
    ind = which(d == min(d))
    c(sensitivity = y[[ind]], specificity = 1-x[[ind]],
      cutoff = p[[ind]])
  }, perf@x.values, perf@y.values, pred@cutoffs)
}
#print(opt.cut(roc.perf.knn, knnpred))
roc.result = as.data.frame((opt.cut(roc.perf.knn, knnpred)))
roc.result
```

```
##           V1
## sensitivity 0.960
## specificity 0.888
## cutoff     0.800
```


```
knn.sens = roc.result[1,]
knn.spec = roc.result[2,]
knn.cutoff = roc.result[3,]

auc.perf.knn = performance(knnpred, measure = 'auc')
auc.knn = auc.perf.knn@y.values
auc.knn
```

```
## [[1]]  
## [1] 0.96
```

# XGBoost

```
#Prepare model for ROC curve  
xgbpred = prediction(xgb.pred.prob, test$y)  
  
roc.perf.xgb = performance(xgbpred, measure = "tpr", x.measure = "fpr")  
plot(roc.perf.xgb, main='ROC XGBoost', colorize=T)
```



```
#Optimal cutoff  
opt.cut = function(perf, pred){  
  cut.ind = mapply(FUN=function(x, y, p){  
    d = (x - 0)^2 + (y-1)^2  
    ind = which(d == min(d))  
    c(sensitivity = y[[ind]], specificity = 1-x[[ind]],  
      cutoff = p[[ind]])  
  }, perf@x.values, perf@y.values, pred@cutoffs)  
}  
#print(opt.cut(roc.perf.xgb, xgbpred))  
roc.result = as.data.frame((opt.cut(roc.perf.xgb, xgbpred)))  
roc.result
```

```
##          V1
## sensitivity 0.915
## specificity 0.883
## cutoff     0.597
```

```
xgb.sens = roc.result[1,]
xgb.spec = roc.result[2,]
xgb.cutoff = roc.result[3,]

auc.perf.xgb = performance(xgbpred, measure = 'auc')
auc.xgb = auc.perf.xgb@y.values
auc.xgb
```


```
## [[1]]
## [1] 0.954
```

## Adaboost

```
#Prepare model for ROC curve
ada.pred <- predict(ada.cls, x.testA, type = "prob")

adapred = prediction(ada.pred, test$y)

roc.perf.ada = performance(adapred, measure = "tpr", x.measure = "fpr")
plot(roc.perf.ada, main='ROC Adaboost', colorize=T)
```





```

#Optimal cutoff
opt.cut = function(perf, pred){
  cut.ind = mapply(FUN=function(x, y, p){
    d = (x - 0)^2 + (y-1)^2
    ind = which(d == min(d))
    c(sensitivity = y[[ind]], specificity = 1-x[[ind]],
      cutoff = p[[ind]])
  }, perf@x.values, perf@y.values, pred@cutoffs)
}
#print(opt.cut(roc.perf.ada, adapred))
roc.result = as.data.frame((opt.cut(roc.perf.ada, adapred)))
roc.result

```

```

##              V1
## sensitivity 0.926
## specificity 0.896
## cutoff      0.605

```

```

ada.sens = roc.result[1,]
ada.spec = roc.result[2,]
ada.cutoff = roc.result[3,]

auc.perf.ada = performance(adapred, measure = 'auc')
auc.ada = auc.perf.ada@y.values
auc.ada

```

```

## [[1]]
## [1] 0.961

```

## Gradient boosting

```


#Prepare model for ROC curve
gbm.pred <- predict (gbm.fit, test, type = "prob")

gbmpred = prediction(gbm.pred[,2], test$y)

roc.perf.gbm = performance(gbmpred, measure = "tpr", x.measure = "fpr")
plot(roc.perf.gbm, main='ROC Gradient boosting', colorize=T)

```

## ROC Gradient boosting



```
#Optimal cutoff
opt.cut = function(perf, pred){
  cut.ind = mapply(FUN=function(x, y, p){
    d = (x - 0)^2 + (y-1)^2
    ind = which(d == min(d))
    c(sensitivity = y[[ind]], specificity = 1-x[[ind]],
      cutoff = p[[ind]])
  }, perf@x.values, perf@y.values, pred@cutoffs)
}
#print(opt.cut(roc.perf.gbm, gbmpred))
roc.result = as.data.frame((opt.cut(roc.perf.gbm, gbmpred)))
roc.result
```

```
##          V1
## sensitivity 0.911
## specificity 0.868
## cutoff     0.563
```

```
gbm.sens = roc.result[1,]
gbm.spec = roc.result[2,]
gbm.cutoff = roc.result[3,]

auc.perf.gbm = performance(gbmpred, measure = 'auc')
auc.gbm = auc.perf.gbm@y.values
auc.gbm
```

```
## [[1]]
## [1] 0.949
```

## AUC Summary

```
options(digits = 3)
cl.err <- matrix(c(test.err.rf, FNR.rf, rf.cutoff, (1-rf.sens), auc.rf,
                  test.err.knn, FNR.knn, knn.cutoff, (1-knn.sens), auc.knn,
                  test.err.xgb, FNR.xgb, xgb.cutoff, (1-xgb.sens), auc.xgb,
                  test.err.ada, FNR.ada, ada.cutoff, (1-ada.sens), auc.ada,
                  test.err.gbm, FNR.gbm, gbm.cutoff, (1-gbm.sens), auc.gbm
                  ),
               ncol=5, byrow=TRUE)
colnames(cl.err) <- c('misclass err', 'Type-II err@0.5', 'cutoff', 'Type-II err@cutoff', 'AUC')
rownames(cl.err) <- c('RandomForest',
                     'KNN',
                     'XGBoost',
                     'Adaboost',
                     'Gradboost')

as.matrix(cl.err)
```

##	misclass err	Type-II err@0.5	cutoff	Type-II err@cutoff	AUC
## RandomForest	0.0412	0.000136	0.784	0.0101	0.998
## KNN	0.0907	0.0127	0.8	0.0398	0.96
## XGBoost	0.0988	0.0533	0.597	0.0854	0.954
## Adaboost	0.0859	0.0477	0.605	0.0737	0.961
## Gradboost	0.111	0.0711	0.563	0.089	0.949